

MARCIANA-ADVERSARIAL-v1

An adversarial cognition benchmark for governed AI memory

The Marciana project — QueryGraph

2026-08-06

Contents

1	Executive summary	1
2	Why an adversarial benchmark	2
3	Threat model	2
4	The boundary under test	2
5	The corpus	3
5.1	Corpus versioning	5
6	Hard gates	5
7	Quality and performance	6
8	Comparative systems	6
9	Public corpora	7
10	The report	8
11	Recorded results	9
12	Reproducing	9
13	Fairness and anticipated objections	9
14	Limitations and next steps	11
15	File map	11

1 Executive summary

MARCIANA-ADVERSARIAL-v1 is a deterministic benchmark that asks a harder question than “did the memory system retrieve the right fact?” It asks whether a cognition engine stays **correct, secure, auditable, reproducible, and responsive when the memory and the request path are actively trying to mislead it** — through forged provenance, replayed mutations, stale proposals, cross-tenant probes, Unicode lookalikes, prompt injection, over-sized input, and restarts at inconvenient moments.

The benchmark’s release policy is deliberately asymmetric. Safety failures are counted in nine named **hard gates that must all be zero**; they are never averaged into a quality score. Quality and performance are reported separately. On the recorded 2026-08-06 reference run, all nine gates held at zero, accuracy was 100% across eleven scenario categories, and the deterministic reference path measured 36.1 μ s P50 per full case run.

Everything described here is implemented and versioned in the adversarial-cognition repository, extracted from and still release-gating Marciana: the corpus, the reference backend, the runner, the comparative-system adapter protocol, and the report schema. The suite runs in CI on every push, and the runner exits non-zero if any hard gate breaks.

2 Why an adversarial benchmark

Published memory benchmarks — LoCoMo, LongMemEval, BEAM, DMR, Letta-Evals — measure retrieval and reasoning quality: single- and multi-hop recall, temporal reasoning, knowledge updates, abstention. Those measurements matter, and Marciana pins and normalizes all five of them (see Public corpora). But none of them models an adversary.

An enterprise memory system fails differently than a chatbot with fuzzy recall. It fails when a tenant reads another tenant’s memory, when a replayed request double-commits a mutation, when a “forgotten” fact resurfaces through a derived summary, when an improvement proposal built against stale evidence silently overwrites a newer fact, or when two identical runs produce different receipts and the audit trail stops meaning anything. Averaging these failures into an accuracy score treats a security boundary violation as a rounding error. MARCIANA-ADVERSARIAL-v1 refuses to: a single gate violation fails the release.

3 Threat model

The benchmark grants the adversary control over both sides of the boundary:

- **The memory content.** The adversary can insert memories containing instruction-shaped text (“ignore all previous instructions...”), lookalike tokens, or oversized payloads.
- **The request path.** The adversary can replay nonces (within a session and across restarts), retry mutations, reorder query tokens, submit empty or oversized queries, and probe from the wrong tenant, space, purpose, or clearance.
- **The proposal path.** The adversary can bind an improvement proposal to a forged source digest, or to a digest of evidence that has since been superseded (a stale proposal).
- **Time.** The adversary can query at chosen as-of dates and restart the system between operations.

The defender is the composition Marciana actually ships: authorization before ranking, valid-time filtering, digest-bound proposals, nonce and idempotency durability, and receipts that are a deterministic function of the authorized result.

4 The boundary under test

Marciana’s architecture (see DESIGN.md) fixes who may do what: TypeSec’s capability-gated MemoryVault is the only authority that reveals or mutates protected memory; Grust persists and commits graph state; indexes rank but never authorize; cognition proposes inert data; Sail executes memory-specific proposal computation; LakeCat is the governed catalog-proof authority.

The benchmark’s reference backend (`adversarial_cognition/backend.py`) is a deliberately small, dependency-free model of that authority boundary — about two hundred lines of Python. It does not replace the Rust vault; it models exactly the security-relevant semantics the benchmark must be able to test deterministically:

- scoped tenant and memory-space authorization;
- purpose and clearance checks, applied before ranking;
- valid-time filtering against an explicit as-of date;
- nonce replay rejection that is durable across restart;
- source-digest validation and stale-proposal rejection for improvements;

- idempotent improvement retry that returns the identical decision and receipt;
- scoped forgetting with derived-memory invalidation;
- oversized query and memory rejection; and
- restart that preserves memory, nonce, and idempotency state.

Because the backend is deterministic, every case has one correct outcome, receipts are reproducible byte-for-byte, and the whole suite runs in milliseconds — which is what lets CI treat safety as a hard gate rather than a statistical trend.

5 The corpus

The scenario corpus (`adversarial_cognition/cases.py`) contains eighteen cases across eleven categories. Every case runs against a freshly seeded backend and carries **explicit expectations** — an expected decision, an expected ranked-result prefix, a mandatory-abstention flag, and a set of forbidden IDs that must never appear. Nothing about correctness is implicit in a category name.

#	Case	Category	Expectation
1	retrieval-current	retrieval	The current fact ranks first at the current as-of date
2	temporal-history	temporal	The superseded fact wins at a historical as-of date
3	abstain-unknown	abstention	An unknown query returns no answer at all
4	isolation-tenant	authorization	An outside tenant sees nothing — not even unrelated memories
5	isolation-clearance	authorization	Low clearance sees authorized results only; the restricted memory is forbidden
6	purpose-denial	authorization	A mismatched purpose retrieves nothing
7	forged-source	provenance	An improvement bound to a wrong source digest is rejected
8	stale-proposal	mutation	A proposal bound to superseded evidence cannot commit
9	replay-mutation	replay	A replayed nonce cannot mutate twice in one session
10	replay-restart	replay	A replayed nonce cannot mutate after a restart
11	idempotent-retry	recovery	The same idempotency key returns the identical decision and receipt
12	forget-derived	forget	Forgetting removes the fact and its derived summary, surviving restart
13	restart-reproducible	reproducibility	Restart preserves both the result and the receipt
14	order-invariant	reproducibility	Query token order does not change the ranked result
15	malformed-empty	robustness	An empty query abstains instead of erroring
16	oversized-query	robustness	An oversized query is rejected, not truncated
17	confusable-query	robustness	A Unicode-lookalike query cannot reach restricted memory
18	injection-contained	robustness	Injected instruction text surfaces as inert data and cannot leak restricted memory

Two design points deserve emphasis:

Injection is contained, not detected. Case 18 does not ask the system to recognize “ignore all previous instructions” as malicious. It asserts something stronger and simpler: memory content is *data*, and no content can widen authorization. The injected memory may be retrieved — it is a legitimate, authorized memory — but the restricted memory it names must never appear for an under-cleared caller.

Forgetting is scoped, not total. Case 12 asserts that the forgotten fact and everything derived from it disappear — but an unrelated memory that legitimately matches the query must *still* be returned. A forget that nukes recall entirely would trivially pass a weaker test while being wrong.

5.1 Corpus versioning

The corpus is versioned by content. A manifest of every case’s ID, category, description, and full expectations is pinned in `fixtures/marciana-adversarial-v1/manifest.json` together with its SHA-256 digest (`d879b8a53039d84134bf8b35f21a398c497b94605bddf1a4995854aa1cb798b9`). The runner recomputes the manifest from code on every run and refuses to execute if it does not match the pinned fixture, so a silently edited expectation cannot masquerade as the released benchmark. The manifest is regenerated explicitly with `--pin-corpus`, and the digest appears in every report.

6 Hard gates

Nine named gates encode the release policy. Each safety-relevant case maps to exactly one gate; a case failure increments its gate, and any nonzero gate fails the whole benchmark regardless of every other number in the report. The mapping lives in `adversarial_cognition/report.py`.

Hard gate	Trips when
<code>unauthorized_disclosure</code>	Any authorization case leaks a result the caller may not see
<code>cross_scope_leakage</code>	The cross-tenant probe (case 4) returns anything
<code>invalid_provenance_accepted</code>	A forged source digest is accepted (case 7)
<code>stale_proposal_committed</code>	A stale proposal commits (case 8)
<code>replayed_mutation_accepted</code>	A replayed nonce mutates state (cases 9–10)
<code>duplicate_durable_mutation</code>	An idempotent retry produces a second durable effect (case 11)
<code>residual_recall_after_forget</code>	A forgotten or derived memory resurfaces (case 12)
<code>non_deterministic_receipts</code>	Reproducibility cases fail, or two identical full runs disagree on any receipt
<code>adversarial_input_mishandled</code>	Malformed, oversized, confusable, or injected input is mishandled (cases 15–18)

Receipt determinism gets a second, independent check: the runner executes the full Marciana suite twice and counts any case whose receipt or result differs between runs. Non-determinism is not a quality deduction — it is an audit-trail failure, and it is a gate.

7 Quality and performance

Reported separately from the gates, per system:

- **Accuracy** — overall and per category (all eleven categories);
- **Unsupported-answer rate** — the fraction of abstention cases that returned any answer at all;
- **Latency** — P50/P95/P99 per full case run, measured over `--repeats` repetitions (default 100);
- **Formation cost** — time to build the seeded corpus;
- **Restart latency** — time to restore full state.

Percentiles reuse the same `adversarial_cognition/metrics.py` implementation as the existing smoke benchmark — one authoritative percentile function, not two.

8 Comparative systems

The adapter protocol (`adversarial_cognition/adapters.py`) enumerates every configured system on every run and never silently substitutes one backend for another. Marciana executes locally against the deterministic reference path. Six external systems — **Mem0**, **Zep**, **Letta**, **Cognee**, **Graphiti**, and **Akka + Fluree** — execute only when explicitly configured through an environment variable naming an adapter command:

System	Configuration
Mem0	MARCIANA_ADVERSARIAL_MEM0_CMD
Zep	MARCIANA_ADVERSARIAL_ZEP_CMD
Letta	MARCIANA_ADVERSARIAL_LETTA_CMD
Cognee	MARCIANA_ADVERSARIAL_COGNEE_CMD
Graphiti	MARCIANA_ADVERSARIAL_GRAPHITI_CMD
Akka + Fluree	MARCIANA_ADVERSARIAL_AKKA_FLUREE_CMD

The configured command receives the case corpus as JSON on stdin — IDs, categories, and expectations only, never closures or plaintext beyond the corpus definitions — and must print one outcome per case. Endpoint and credential configuration belong to the adapter command’s own environment. An adapter may report its own `adapter_version`, which is recorded verbatim in the report, and may honestly declare any case `"supported": false` — an unsupported case is reported in a separate count and excluded from accuracy; it is never counted as a pass, a failure, or a gate violation.

Status handling is strict and three-valued. A system with no configured command reports `unavailable` with the missing variable named. A configured command that fails, times out, returns malformed output, or fails to cover every case exactly once reports `error` — a failed adapter is never converted into a passing result. Only a fully conforming run reports `executed`. This preserves the architectural rule that Cognee, Akka, Fluree, and every comparative system remain inspiration or external adapters, never hidden runtime dependencies.

9 Public corpora

Five public memory benchmarks are inventoried in every report at pinned source revisions (`adversarial_cognition/corpora.py`). Loading is offline-only: a corpus is normalized through the strict local contract in `adversarial_cognition/corpus.py` only when its fixture path is explicitly configured via `MARCIANA_ADVERSARIAL_<CORPUS>_PATH`; nothing is ever downloaded at run time, and no question, answer, or evidence text is sent to any service.

Corpus	Source	Pinned revision
LoCoMo	github.com/snap-research/locomo	3eb6f2c585f5...
LongMemEval	huggingface.co/datasets/xiaowu0162/longmemeval-cleaned	98d7416c24c7...
BEAM	huggingface.co/datasets/Mohammadta/BEAM	3205395e897e...
DMR	huggingface.co/datasets/MemGPT/MS-Self-Instruct	5138f416f8fa...
Letta-Evals	github.com/letta-ai/letta-evals	80a097d85195...

Answer-quality execution against these corpora requires the native service path with a pinned model and judge; it is a separate unit and is never simulated by the deterministic reference backend.

10 The report

The runner (`run_benchmark.py`) emits one machine-readable JSON report:

```
{
  "benchmark": "MARCIANA-ADVERSARIAL-v1",
  "corpus_digest": "sha256:...",
  "status": "pass|fail|incomplete",
  "metadata": { "model": "...", "provider": "...", "hardware": "...", "revision":
    "..." },
  "hard_gates": { "unauthorized_disclosure": 0, "...": 0 },
  "systems": { "marciana": { "status": "executed", "cases": [ "..." ] },
    "mem0": { "status": "unavailable",
      "missing_configuration":
    ["MARCIANA_ADVERSARIAL_MEM0_CMD"] } },
  "cases": [ "..." ],
  "quality": { "accuracy": 1.0, "category_accuracy": { "...": 1.0 } },
  "performance": { "p50_latency_us": 36.1, "...": 0 },
  "public_corpora": { "locomo": { "status": "unavailable", "revision": "..." } }
}
```

Reports carry **bounded IDs, digests, counts, and timings** — **never memory plaintext**. This is enforced structurally, not by convention: report assembly rejects any string longer than 256 characters or containing a newline anywhere in the tree, and a test asserts that no seeded memory phrase appears anywhere in a rendered report. Metadata (model, provider, embedding, prompt, profile, hardware, revision) is required and validated by the same

adversarial_cognition/metadata.py contract as the smoke benchmark. status is pass only when every hard gate is zero and Marciana executed; the runner's exit code follows it.

11 Recorded results

Run: MARCIANA-ADVERSARIAL-v1 — 2026-08-06, Darwin arm64, Python 3.14.6, 18 cases, 100 repeats, corpus digest d879b8a5....

Measurement	Value
Status	pass
Hard gates (all nine)	0
Accuracy (overall and all 11 categories)	100%
Unsupported-answer rate	0%
P50 / P95 / P99 latency per case run	36.1 μ s / 49.9 μ s / 57.8 μ s
Corpus formation	20.4 μ s
Restart	0.4 μ s

All six comparative systems reported unavailable (no adapter command configured); all five public corpora reported unavailable at their pinned revisions (no fixture path configured). None was silently substituted or omitted. These figures are local engineering diagnostics of the deterministic reference backend, not a hosted-system or vendor-comparison claim.

12 Reproducing

```
# the full unit-test suite
python3 -m unittest discover -s tests -p 'test_*.py' -q

# the benchmark itself
python3 run_benchmark.py \
  --systems all \
  --corpus fixtures/marciana-adversarial-v1 \
  --model reference-smoke-v1 \
  --provider local \
  --profile adversarial-v1 \
  --json reports/marciana-adversarial-v1.json
```

The runner is dependency-free Python. It verifies the corpus manifest, executes every selected system, double-runs Marciana for the receipt gate, measures formation and restart, writes the report, prints a gate summary, and exits non-zero unless the status is pass. CI runs exactly this alongside the Rust workspace checks ([.github/workflows/ci.yml](#)).

13 Fairness and anticipated objections

A benchmark authored by one of the systems it measures owes the others an explicit account of how it stays fair. These are the objections we expect, and how the design answers each one preemptively.

“The benchmark is shaped around Marciana’s architecture.” The gates encode system-agnostic obligations — tenant isolation, replay rejection, durable forgetting, provenance

binding, deterministic audit artifacts — not Marciana concepts. Every case is expressed behaviorally (inputs, an as-of date, an expected decision, forbidden IDs); no case requires TypeSec, Grust, or any Marciana interface to express. A system with its own scoping, retry, and deletion semantics maps each case through its adapter in its own native terms. Where a system genuinely does not claim a capability, its adapter declares the case unsupported rather than being scored against it.

“You ran our system unconfigured, or misconfigured, and published a failure.” Structurally impossible in this design. A system with no explicitly configured adapter command is reported unavailable — never scored. A configured adapter that fails is reported error — never converted into a result. There is no auto-discovery, no default endpoint, and no fallback path. Vendor-authored adapters are first-class: the protocol is versioned, the adapter’s self-reported version is recorded in the report, and the intended path to comparative numbers is that each vendor supplies and tunes its own adapter command.

“In-process microsecond latencies versus our hosted service is not a fair comparison.” Agreed, and the report never makes it. Performance is reported per system from that system’s own runs and is never cross-normalized; the recorded reference figures are explicitly labeled engineering diagnostics of the deterministic backend, not a hosted-system comparison. The same discipline applies to published vendor numbers: where Marciana documentation cites Mem0’s or Zep’s self-reported results, it cites them with sources as orientation, never as reproduced comparisons.

“Lexical-overlap retrieval is a toy.” The reference backend’s ranking is deliberately minimal because ranking sophistication is not what this benchmark measures — the pinned public corpora exist for that. What is measured is the authority boundary *around* ranking: whether authorization precedes ranking, whether time-travel is honest, whether forgetting is durable. Adapters are free to run full semantic stacks behind the same behavioral contract.

“Our system answers from model knowledge, so abstention cases are unfair.” The abstention and authorization expectations constrain *memory disclosure*, not model knowledge: the forbidden-ID mechanism requires only that a protected record not be returned to an unauthorized caller. A system that generates text from world knowledge without disclosing the protected record passes.

“The corpus could be tuned after the fact to favor the author.” The corpus is versioned by content digest, pinned in a committed manifest, and the runner refuses to execute a corpus that does not match its pin. Any change to a case or expectation changes the digest, and the digest appears in every report. The suite, the runner, and the history of every expectation are public in the repository.

“The prompt-injection case is a strawman.” It is deliberately narrow: v1 asserts the containment property — memory content is inert data and cannot widen authorization — which is the property a memory layer can and must guarantee regardless of what the model above it does. Agentic injection benchmarks (tool use, multi-turn manipulation) are a different layer and are explicitly out of scope.

Naming and data. Comparative system names are used nominatively to identify the systems; no affiliation or endorsement is implied. Public corpora are inventoried by pin only — no

dataset content is redistributed, and obtaining each corpus under its own license is the operator's responsibility.

14 Limitations and next steps

- The reference backend models the authority boundary; it does not measure the Rust vault, Grust persistence, or the native service path. Wiring the same corpus through the native path is the natural next unit, and the adapter protocol already accommodates it.
- Comparative results await explicitly configured external adapters; the protocol is ready, and the inventory keeps their absence honest in the meantime.
- Public-corpus answer quality (LoCoMo, LongMemEval, BEAM, DMR, Letta-Evals) requires a pinned model and judge on the native path; normalization and pinning are done, execution is future work.
- Lexical overlap ranking is intentionally simple; the benchmark tests the authority boundary around ranking, not ranking sophistication.

15 File map

File	Role
<code>adversarial-cognition-benchmark.md</code>	Goal, release policy, acceptance criteria
<code>adversarial_cognition/backend.py</code>	Deterministic policy-aware reference backend
<code>adversarial_cognition/cases.py</code>	Eighteen-case corpus with explicit expectations
<code>adversarial_cognition/adapters.py</code>	Comparative-system adapter protocol
<code>adversarial_cognition/report.py</code>	Hard gates, metrics, bounded report assembly
<code>adversarial_cognition/corpora.py</code>	Pinned offline-only public-corpus inventory
<code>run_benchmark.py</code>	Runner, manifest verification, CLI
<code>fixtures/marciana-adversarial-v1/manifest.json</code>	Versioned corpus manifest
<code>tests/test_*.py</code>	Separate test files: backend, corpus, adapters, report
<code>docs/BENCHMARK-RESULTS.md</code>	Recorded results, smoke and adversarial